

#### **CUDA**



- Programing system for machines with GPUs
  - Programming Language
  - Compilers
  - Runtime Environments
  - Drivers
  - Hardware





# **CUDA: Heterogeneous Parallel Computing**



- CPU optimized for fast single-thread execution
  - Cores designed to execute 1 thread or 2 threads concurrently
  - Large caches attempt to hide DRAM access times
  - Cores optimized for low-latency cache accesses
  - Complex control-logic for speculation and out-of-order execution

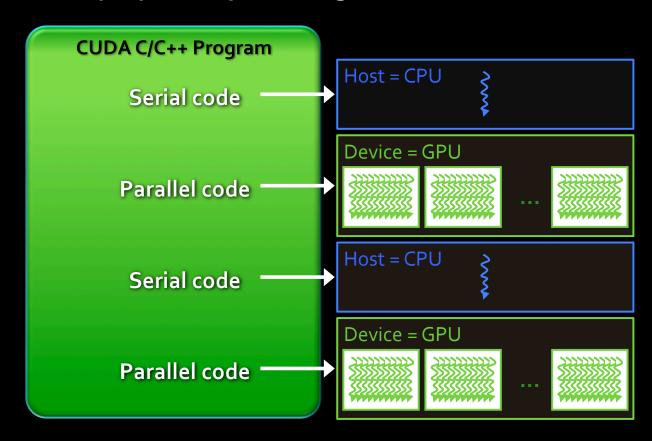


- GPU optimized for high multi-thread throughput
  - Cores designed to execute many parallel threads concurrently
  - Cores optimized for data-parallel, throughput computation
  - Chips use extensive multithreading to tolerate DRAM access times

## Anatomy of a CUDA C/C++ Program



- Serial code executes in a Host (CPU) thread
- Parallel code executes in many concurrent Device (GPU) threads across multiple parallel processing elements



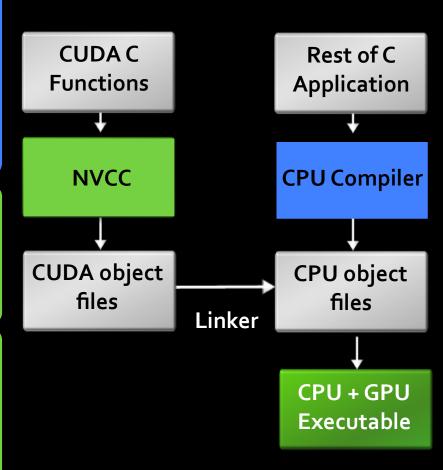
### Compiling CUDA C/C++ Programs



```
// foo.cpp
int foo(int x)
{
    ...
}
float bar(float x)
{
    ...
}
```

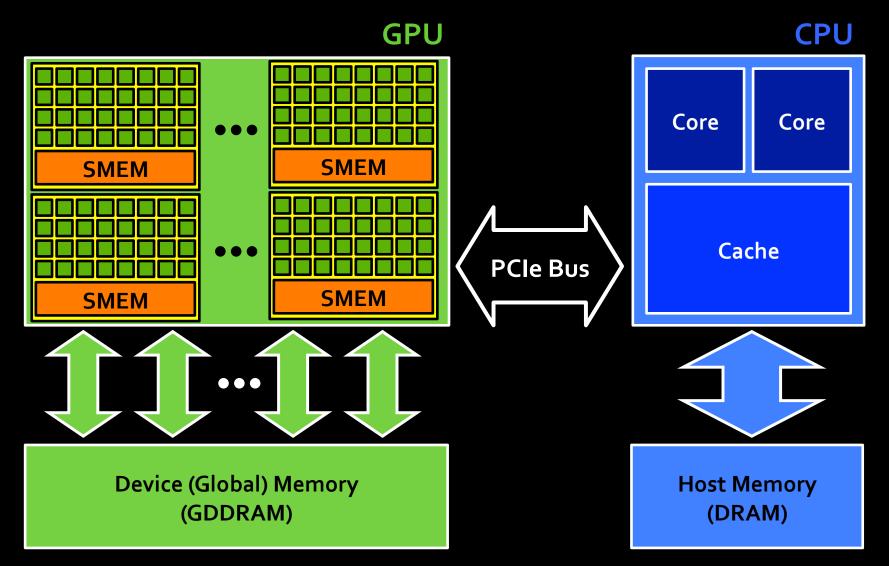
```
// saxpy.cu
__global__ void saxpy(int n, float ... )
{
   int i = threadIdx.x + ...;
   if (i < n) y[i] = a*x[i] + y[i];
}</pre>
```

```
// main.cpp
void main( ) {
  float x = bar(1.0)
  if (x<2.0f)
    saxpy<<<...>>>(foo(1), ...);
  ...
}
```



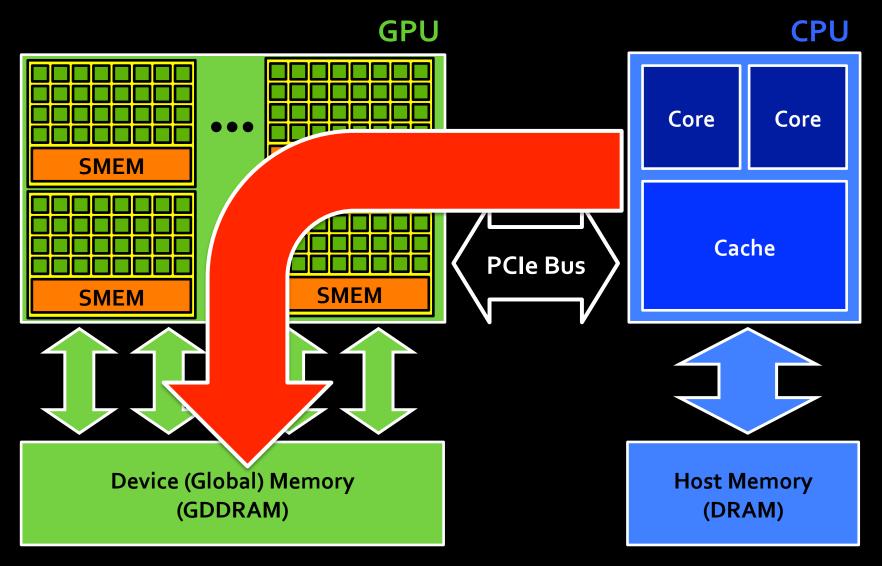
#### Canonical execution flow





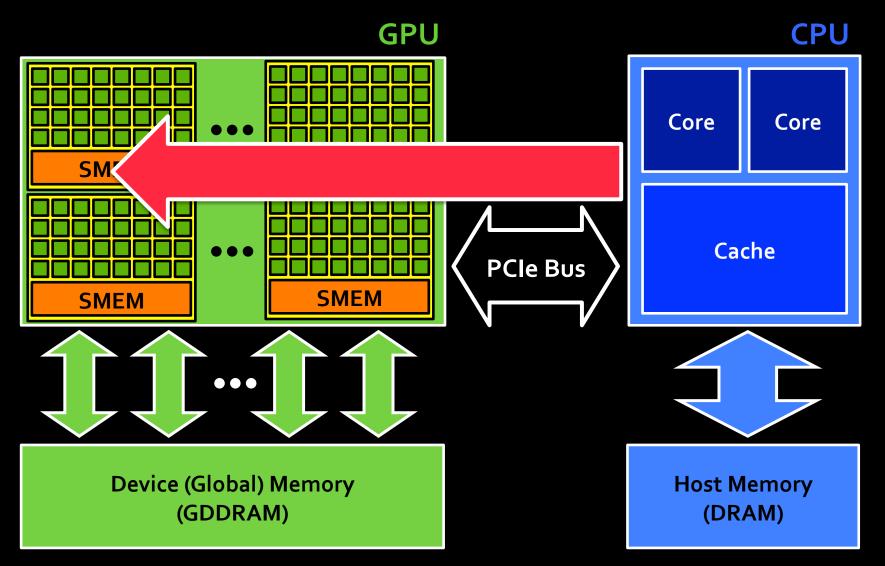
# Step 1 – copy data to GPU memory





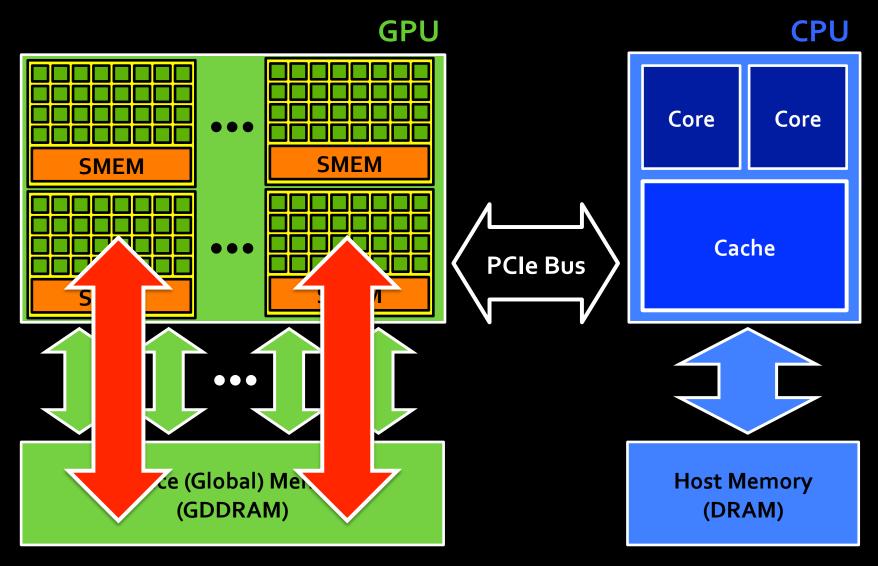
# Step 2 – launch kernel on GPU





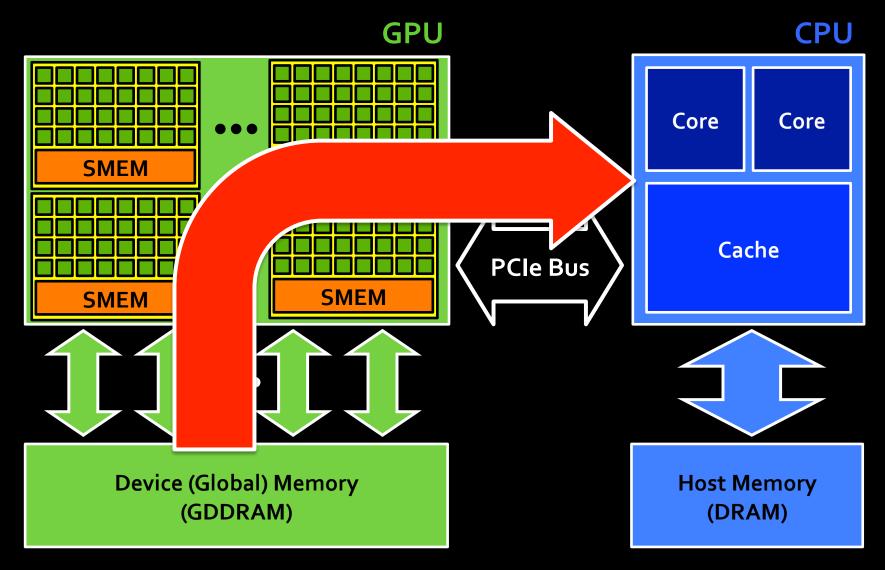
## Step 3 – execute kernel on GPU





# Step 4 – copy data to CPU memory







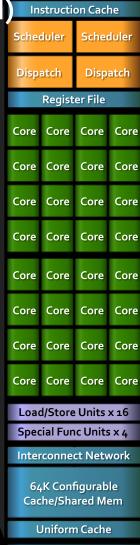
# **CUDA ARCHITECTURE**

#### Contemporary (Fermi) GPU Architecture



32 CUDA Cores per Streaming Multiprocessor (SM)

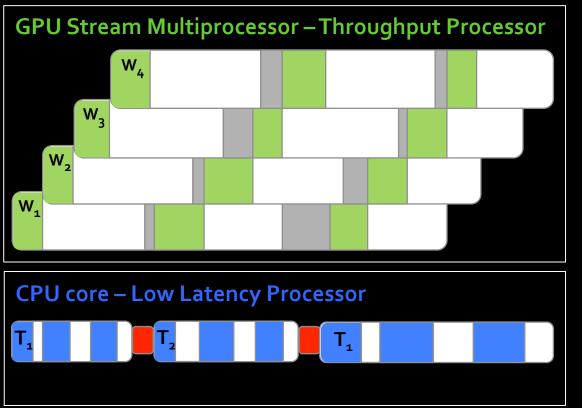
- 32 fp32 ops/clock
- **16 fp64 ops/clock**
- 32 int32 ops/clock
- 2 Warp schedulers per SM
  - 1,536 concurrent threads
- 4 special-function units
- 64KB shared memory + L1 cache
- 32K 32-bit registers
- Fermi GPUs have as many as 16 SMs
  - 24,576 concurrent threads

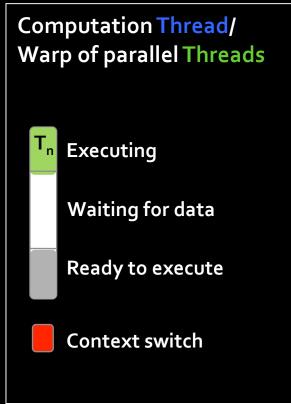


### Multithreading



- CPU architecture attempts to minimize latency within each thread
- GPU architecture hides latency with computation from other thread warps







# **CUDA PROGRAMMING MODEL**

#### **CUDA Kernels**



- Parallel portion of application: execute as a kernel
  - Entire GPU executes kernel
  - Kernel launches create thousands of CUDA threads efficiently

CPU	Host	Executes functions
GPU	Device	Executes kernels

- CUDA threads
  - Lightweight
  - Fast switching
  - 1000s execute simultaneously
- Kernel launches create hierarchical groups of threads
  - Threads are grouped into Blocks, and Blocks into Grids
  - Threads and Blocks represent different levels of parallelism

### **CUDA C: C with a few keywords**



- Kernel: function that executes on device (GPU) and can be called from host (CPU)
  - Can only access GPU memory
  - No variable number of arguments
  - No static variables
- Functions must be declared with a qualifier

```
__global__ : GPU kernel function launched by CPU, must return void
__device__ : can be called from GPU functions
__host__ : can be called from CPU functions (default)
__host__ and __device__ qualifiers can be combined
```

- Qualifiers determines how functions are compiled
  - Controls which compilers are used to compile functions

#### **CUDA Kernels: Parallel Threads**



- A kernel is a function executed on the GPU as an array of parallel threads
- All threads execute the same kernel code, but can take different paths
- Each thread has an ID
  - Select input/output data
  - Control decisions

```
in[i] in[i+1] in[i+2] in[i+3]

float x = in[threadIdx.x];
float y = func(x);
out[threadIdx.x] = y;

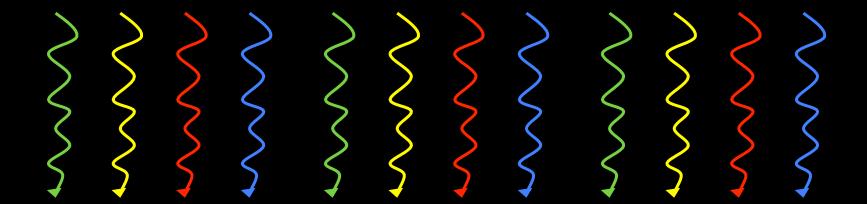
out[i] out[i+1] out[i+2] out[i+3]
```



# **CUDA THREADS**

#### **CUDA Thread Organization**

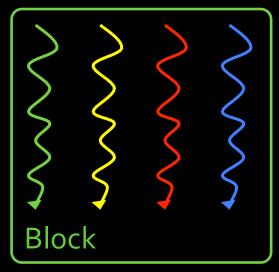


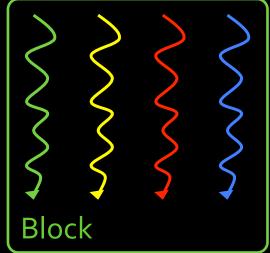


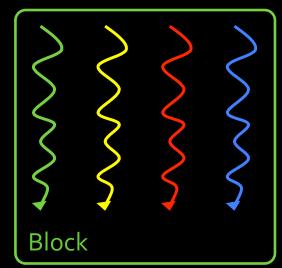
- GPUs can handle thousands of concurrent threads
- CUDA programming model supports even more
  - Allows a kernel launch to specify more threads than the GPU can execute concurrently
  - Helps to amortize kernel launch times

#### **Blocks of threads**





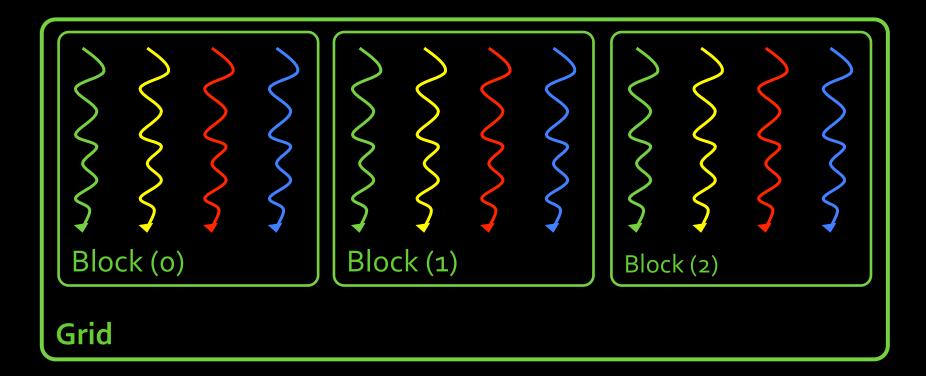




Threads are grouped into blocks

#### **Grids of blocks**





- Threads are grouped into blocks
- Blocks are grouped into a grid
- A kernel is executed as a grid of blocks of threads

#### **Blocks execute on Streaming Multiprocessors**

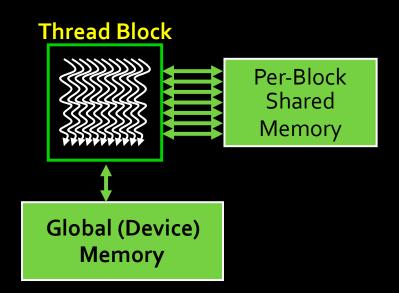


#### **Streaming Processor**

#### **Streaming Multiprocessor**



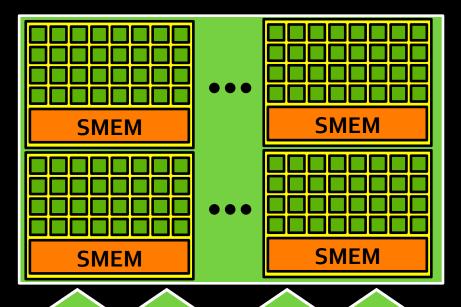
# Registers Global (Device) Memory



#### Grids of blocks executes across GPU



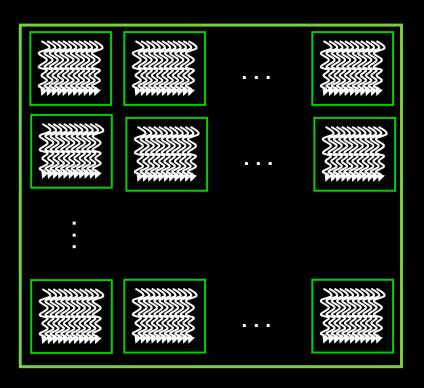
#### **GPU**



Global (Device) Memory

 $\bullet \bullet \bullet$ 

#### **Grid of Blocks**



#### **Kernel Execution Recap**

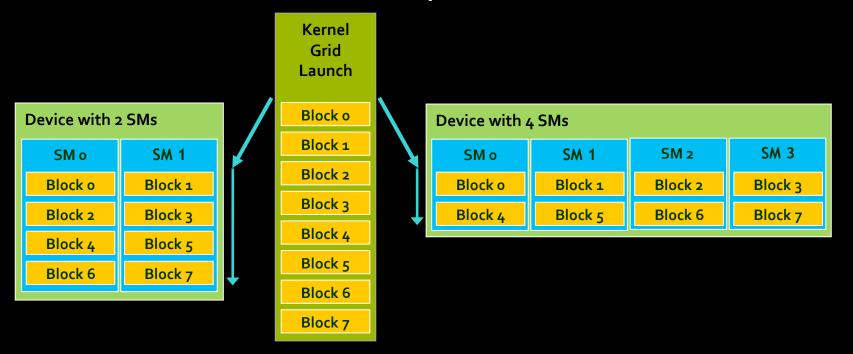


- A thread executes on a single streaming processor
  - Allows use of familiar scalar code within kernel
- A block executes on a single streaming multiprocessor
  - Threads and blocks do not migrate to different SMs
  - All threads within block execute in concurrently, in parallel
- A streaming multiprocessor may execute multiple blocks
  - Must be able to satisfy aggregate register and memory demands
- A grid executes on a single device (GPU)
  - Blocks from the same grid may execute concurrently or serially\*
  - Blocks from multiple grids may execute concurrently
  - A device can execute multiple kernels concurrently





- Blocks may execute in arbitrary order, concurrently or sequentially, and parallelism increases with resources
  - Depends on when execution resources become available
- Independent execution of blocks provides scalability
  - Blocks can be distributed across any number of SMs



#### Blocks also enable efficient collaboration



- Threads often need to collaborate
  - Cooperatively load/store common data sets
  - Share results or cooperate to produce a single result
  - Synchronize with each other
- Threads in the same block
  - Can communicate through shared and global memory
  - Can synchronize using fast synchronization hardware
- Threads in different blocks of the same grid
  - Cannot synchronize reliably
  - No guarantee that both threads are alive at the same time

## Blocks must be independent

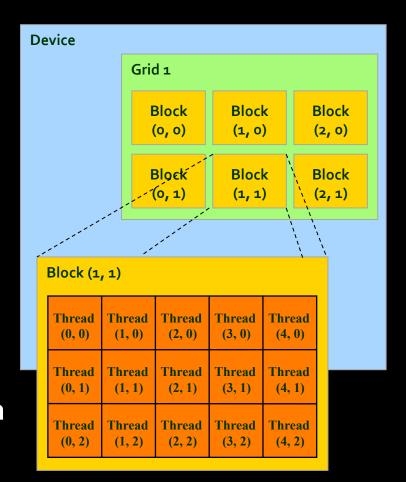


- Any possible interleaving of blocks is allowed
  - Blocks presumed to run to completion without pre-emption
  - May run in any order, concurrently or sequentially
- Programs that depend on block execution order within grid for correctness are not well formed
  - May deadlock or return incorrect results
- Blocks may coordinate but not synchronize
  - shared queue pointer: OK
  - shared lock: **BAD** ... can easily deadlock

#### Thread and Block ID and Dimensions



- Threads
  - 3D IDs, unique within a block
- Thread Blocks
  - 2D IDs, unique within a grid
- Dimensions set at launch
  - Can be unique for each grid
- Built-in variables
  - threadIdx, blockIdx
  - blockDim, gridDim
- Programmers usually select dimensions that simplify the mapping of the application data to CUDA threads



## **Examples of Indexes and Indexing**



```
_global___ void kernel( int *a )
 int idx = blockIdx.x*blockDim.x + threadIdx.x;
 a[idx] = 7;
                          global void kernel( int *a )
 int idx = blockIdx.x*blockDim.x + threadIdx.x;
 a[idx] = blockIdx.x;
                          Output: 0 0 0 0 1 1 1 1 2 2 2 2 3 3 3 3
global void kernel( int *a )
 int idx = blockIdx.x*blockDim.x + threadIdx.x;
 a[idx] = threadIdx.x;
                          Output: 0 1 2 3 0 1 2 3 0 1 2 3 0 1 2 3
```





```
__global__ void kernel(int *a, int dimx, int dimy)
{
  int ix = blockIdx.x*blockDim.x + threadIdx.x;
  int iy = blockIdx.y*blockDim.y + threadIdx.y;
  int idx = iy * dimx + ix;

a[idx] = a[idx]+1;
}
```

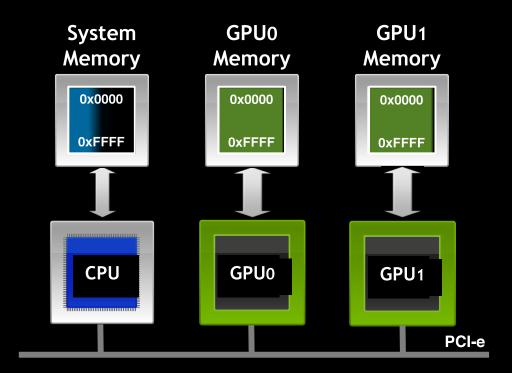


# **CUDA MEMORY MODEL**

#### Independent address spaces

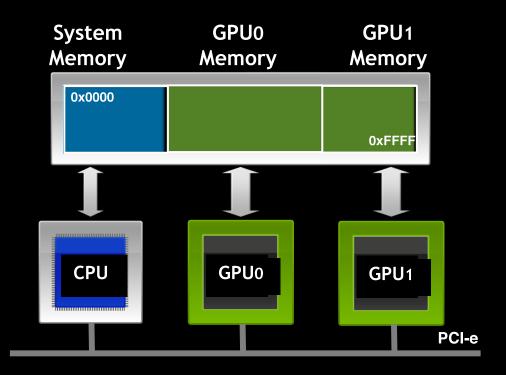


- CPU and GPU have independent memory systems
  - PCIe bus transfers data between CPU and GPU memory systems
- Typically, CPU thread and GPU threads access what are logically different, independent virtual address spaces



# Independent address spaces: consequences

- Cannot reliably determine whether a pointer references a host (CPU) or device (GPU) address from the pointer value
  - Dereferencing CPU/GPU pointer on GPU/CPU will likely cause crash
- Unified virtual addressing (UVA) in CUDA 4.0
  - One virtual address space shared by CPU thread and GPU threads

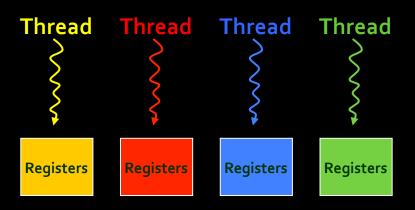


33

# **CUDA Memory Hierarchy**



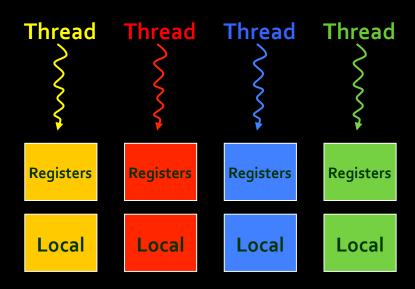
- Thread
  - Registers



# **CUDA Memory Hierarchy**



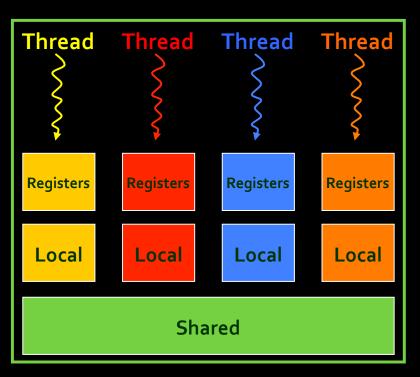
- Thread
  - Registers
  - Local memory



## **CUDA Memory Hierarchy**



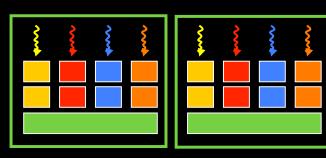
- Thread
  - Registers
  - Local memory
- Thread Block
  - Shared memory



# **CUDA Memory Hierarchy**



- Thread
  - Registers
  - Local memory
- Thread Block
  - Shared memory
- All Thread Blocks
  - Global Memory



Global Memory (DRAM)

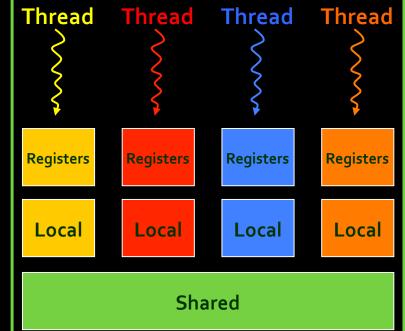
## **Shared Memory**



- Allocated per thread block
- Scope: threads in block
- Data lifetime: same as block
- Capacity: small (about 48kB)
- Latency: a few cycles
- Bandwidth: very high

SM: 32 \* 4 B \* 1.15 GHz / 2 = 73.6 GB/s

GPU: 14 \* 32 \* 4 B \* 1.15 GHz / 2 = 1.03 TB/s

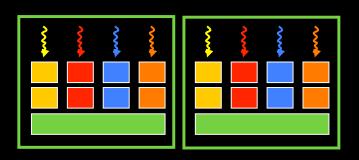


- Common uses
  - Sharing data among threads in a block
  - User-managed cache (to reduce global memory accesses)

## **Global Memory**



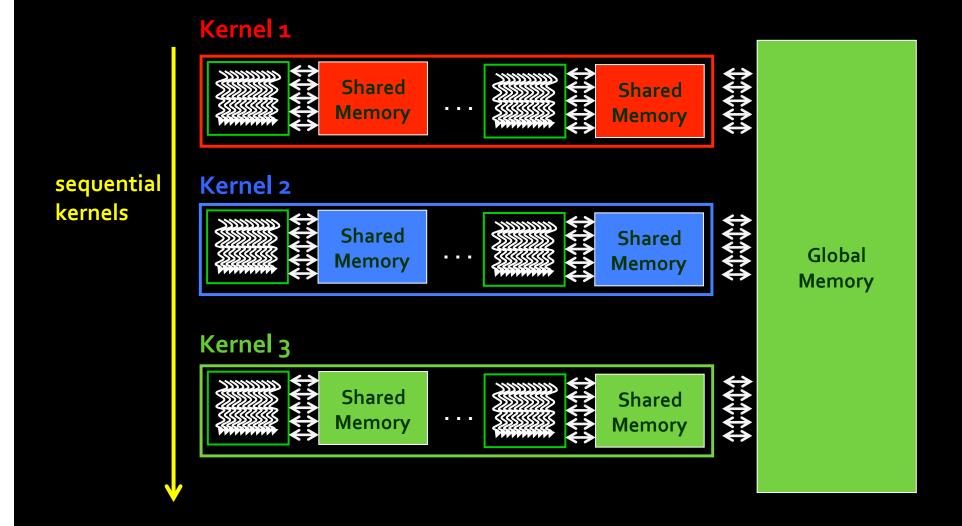
- Allocated explicitly by host (CPU) thread
- Scope: all threads of all kernels
- Data lifetime: determine by host (CPU) thread
  - cudaMalloc (void \*\* pointer, size\_t nbytes)
  - cudaFree (void\* pointer)
- Capacity: large (1-6GB)
- Latency: 400-800 cycles
- Bandwidth: 156 GB/s
  - Data access patterns will limit bandwidth achieved in practice
- Common uses
  - Staging data transfers to/from CPU
  - Staging data between kernel launches



Global Memory (DRAM)

#### **Communication and Data Persistence**





## Managing Device (GPU) Memory



- Host (CPU) manages device (GPU) memory
  - cudaMalloc (void \*\* pointer, size\_t num\_bytes)
  - cudaMemset (void\* pointer, int value, size\_t count)
  - cudaFree(void\* pointer)
- Example: allocate and initialize array of 1024 ints on device

```
// allocate and initialize int x[1024] on device
int n = 1024;
int num_bytes = 1024*sizeof(int);
int* d_x = 0; // holds device pointer
cudaMalloc((void**)&d_x, num_bytes);
cudaMemset(d_x, 0, num_bytes);
cudaFree(d_x);
```

## **Transferring Data**



cudaMemcpy(void\* dst, void\* src, size\_t num\_bytes, enum cudaMemcpyKind

#### direction);

- returns to host thread after the copy completes
- blocks CPU thread until all bytes have been copied
- doesn't start copying until previous CUDA calls complete
- Direction controlled by enum cudaMemcpyKind
  - cudaMemcpyHostToDevice
  - cudaMemcpyDeviceToHost
  - cudaMemcpyDeviceToDevice
- CUDA also provides non-blocking
  - Allows program to overlap data transfer with concurrent computation on host and device
  - Need to ensure that source locations are stable and destination locations are not accessed



# **CUDA EXAMPLE**



```
// [compute] for (i=0; i < n; i++) y[i] = a * x[i] + y[i];
// Each thread processes one element
 global__ void saxpy(int n, float a, float* x, float* y)
 int i = threadIdx.x + blockDim.x * blockIdx.x;
 if (i < n) y[i] = a*x[i] + y[i];
int main()
 // invoke parallel SAXPY kernel with 256 threads / block
 int nblocks = (n + 255)/256;
 saxpy<<<nblocks, 256>>>(n, 2.0, d_x, d_y);
```



```
// [computes] for (i=0; i < n; i++) y[i] = a * x[i] + y[i];
     Each thread processes one element
  global__ void saxpy(int n, float a, float* x, float* y)
 int i = threadIdx.x + blockDim.x * blockIdx.x;
 if (i < n) y[i] = a*x[i] + y[i];
                                                     Device Code
int main()
 // invoke parallel SAXPY kernel with 256 threads / block
 int nblocks = (n + 255)/256;
 saxpy<<<nblocks, 256>>>(n, 2.0, d_x, d_y);
```



```
// [computes] for (i=0; i < n; i++) y[i] = a * x[i] + y[i];
// Each thread processes one element
__global__ void saxpy(int n, float a, float* x, float* y)
 int i = threadIdx.x + blockDim.x * blockIdx.x;
 if (i < n) y[i] = a*x[i] + y[i];
                                                      Host Code
int main()
 int nblocks = (n + 255)/256;
 saxpy<<<nblocks, 256>>>(n, 2.0, d_x, d_y);
```



```
int main()
 // allocate and initialize host (CPU) memory
 float* x = ...;
 float* y = ...;
 // allocate device (GPU) memory
 float *d x, *d y;
 cudaMalloc((void**) &d_x, n * sizeof(float));
 cudaMalloc((void**) &d_y, n * sizeof(float));
 // copy x and y from host memory to device memory
 cudaMemcpy(d_x, x, n*sizeof(float), cudaMemcpyHostToDevice);
 cudaMemcpy(d_y, y, n*sizeof(float), cudaMemcpyHostToDevice);
 // invoke parallel SAXPY kernel with 256 threads / block
 int nblocks = (n + 255)/256;
 saxpy<<<nblocks, 256>>>(n, 2.0, d_x, d_y);
```



```
int main()
 // allocate and initialize host (CPU) memory
 float* x = ...;
 float* y = ...;
 // allocate device (GPU) memory
 float *d x, *d y;
  cudaMalloc((void**) &d_x, n * sizeof(float));
 cudaMalloc((void**) &d_y, n * sizeof(float));
 // copy x and y from host memory to device memory
 cudaMemcpy(d_x, x, n*sizeof(float), cudaMemcpyHostToDevice);
 cudaMemcpy(d_y, y, n*sizeof(float), cudaMemcpyHostToDevice);
  // invoke parallel SAXPY kernel with 256 threads / block
 int nblocks = (n + 255)/256;
  saxpy<<<nblocks, 256>>>(n, 2.0, d_x, d_y);
```



```
// invoke parallel SAXPY kernel with 256 threads / block
int nblocks = (n + 255)/256;
saxpy<<<nblocks, 256>>>(n, 2.0, d_x, d_y);
// copy y from device (GPU) memory to host (CPU) memory
cudaMemcpy(y, d_y, n*sizeof(float), cudaMemcpyDeviceToHost);
// do something with the result...
// free device (GPU) memory
cudaFree(d_x);
cudaFree(d_y;
return 0;
```



```
// invoke parallel SAXPY kernel with 256 threads / block
int nblocks = (n + 255)/256;
saxpy<<<nblocks, 256>>>(n, 2.0, d_x, d_y);
// copy y from device (GPU) memory to host (CPU) memory
cudaMemcpy(y, d_y, n*sizeof(float), cudaMemcpyDeviceToHost);
// do something with the result...
// free device (GPU) memory
cudaFree(d_x);
cudaFree(d_y;
return 0;
```



```
void saxpy_serial(int n, float a, float* x, float* y)
{
  for (int i = 0; i < n; ++i)
    y[i] = a*x[i] + y[i];
}
// invoke host SAXPY function
saxpy_serial(n, 2.0, x, y);
Standard C Code</pre>
```

```
__global__ void saxpy(int n, float a, float* x, float* y)
{
   int i = blockIdx.x*blockDim.x + threadIdx.x;
   if (i < n) y[i] = a*x[i] + y[i];
}
// invoke parallel SAXPY kernel with 256 threads/block
int nblocks = (n + 255) / 256;
saxpy<<<nblocks, 256>>>(n, 2.0, x, y); CUDA C Code
```



# CUDA DEVELOPMENT RESOURCES

# **CUDA Programming Resources**



- CUDA Toolkit
  - Compiler, libraries, and documentation
  - Free download for Windows, Linux, and MacOS
- GPU Computing SDK
  - Code samples
  - Whitepapers
- Instructional materials on NVIDIA Developer site
  - CUDA introduction & optimization webinar: slides and audio
  - Parallel programming course at University of Illinois UC
  - Tutorials
  - Forums

#### **GPU Tools**

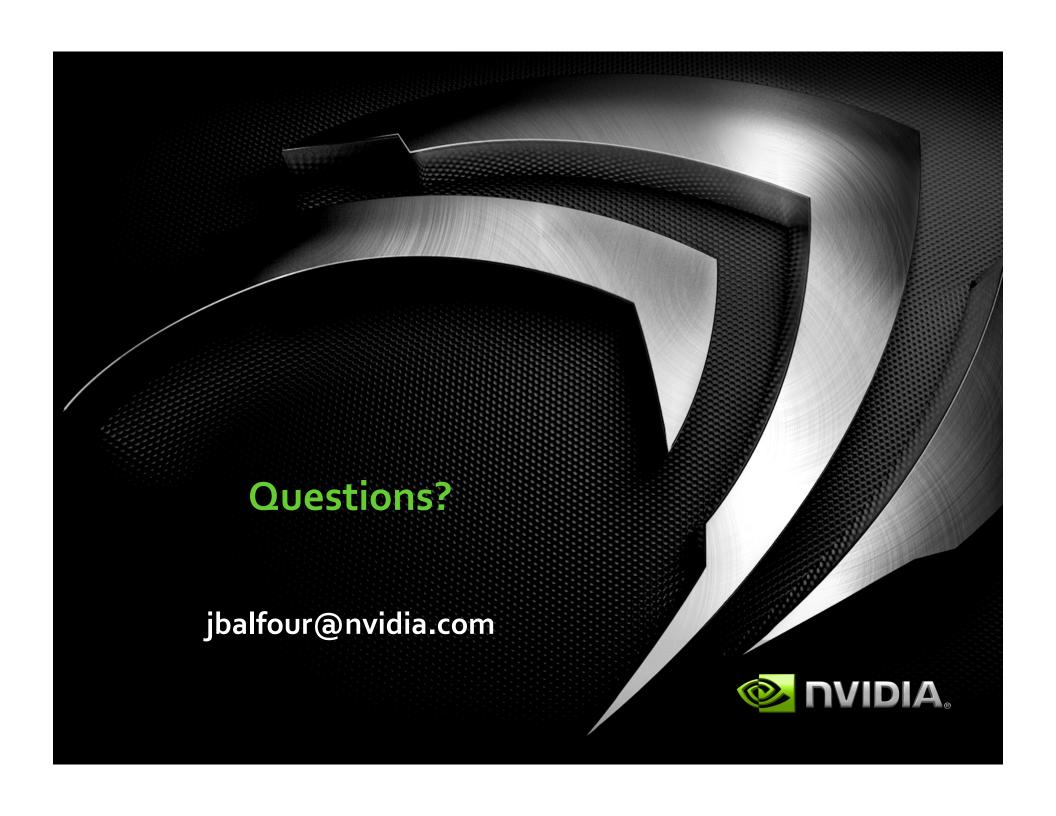


#### Profiler

- Available for all supported OSs
- Command-line or GUI
- Sampling signals on GPU for
  - Memory access parameters
  - > Execution (serialization, divergence)

#### Debugger

- Linux: cuda-gdb
- Windows: Parallel Nsight
- Runs on the GPU



# Acknowledgements



Some slides derived from decks provided by Jared Hoberock and Cliff Woolley